

Lab_14_USA

Team USA: Yuting Deng, Kendra Gilbertson, Lily Durkee, Bennett Hardy

Hermione.Deng@colostate.edu, kendra01@colostate.edu, L.Durkee@colostate.edu, Bennett.Hardy@colostate.edu

2022-11-30

Motivation

The Eurasian lynx (*Lynx lynx*) is a medium-sized predator with broad distribution in the boreal forests of Europe and Siberia. The lynx is classified as a threatened species throughout much of its range and there is controversy about the legal harvest of lynx in Sweden. Proponents of harvest argue that allowing hunting of lynx reduces illegal kill (poaching). Moreover, Sweden is committed to regulate lynx numbers to prevent excessive predation on reindeer because reindeer are critical to the livelihoods of indigenous pastoralists, the Sami. Many environmentalists oppose harvest, however, arguing that lynx are too rare to remove their fully protected status. A similar controversy surrounds management of wolves in the Western United States.

A forecasting model for the abundance of lynx helps managers make decisions that can be justified to citizens. The model you will develop today is not a toy. It is currently used in Sweden and Norway to manage Lynx (H. Andren, N. T. Hobbs, M. Aronsson, H. Broseth, G. Chapron, J. D. C. Linnell, J. Odden, J. Persson, and E. B. Nilsen. Harvest models of small populations of a large carnivore using Bayesian forecasting. Ecological Applications, 30(3):e02063, 2020.)

You have data on the number of lynx family groups censused in a management unit as well as annual records of lynx harvested from the unit. You will model the population using the deterministic model:

$$N_t = \lambda(N_{t-1} - H_{t-1})$$

.

where N_t is the true, unobserved abundance of lynx and H_{t-1} is the number of lynx harvested during $t - 1$ to t . The parentheses in this expression reflect the fact that harvest occurs immediately after census, such that the next years population increment comes from the post-harvest population size.

ADVANCED (for the population modelers) What would be the model if harvest occurred immediately before census? Three months after census? Continuously throughout the year?

Assume the harvest (H_t) is and the number of family groups (y_t) are observed without error. Harvest is closely regulated and all hunters who harvest a lynx are required by law to register the animal with the county. You are entitled to make the assumption that family groups are observed without error because your Scandinavian colleagues are amazing snow trackers and do a good job of estimating the number of family groups (if not the number of lynx) in a management region. The challenge in this problem is that the observations of lynx abundance (family groups) are not the same as the observation of harvest (number of lynx). Fortunately, you have prior information, hard won from radio-telemetry, on the proportional relationship between number of family groups and number of lynx in the population, i.e:

$$\phi = f/N$$

,
where f is the number of family groups and N is the population size, mean $\phi = 0.163$ with standard deviation of the mean = 0.012.

R libraries needed for this lab

You need to load the following libraries. Set the seed to 10 to compare your answers to ours. The data for this problem is located in the LynxFamilies data frame of the BayesNSF package.

```
1 library(BayesNSF)
2 library(rjags)
3 library(MCMCvis)
4 library(HDInterval)
5 set.seed(10)
```

Generating an Informed Prior for ϕ

We've provided you with a useful moment matching function below for converting the mean and standard deviation of ϕ to the parameters for the beta distribution you will use as an informed prior on ϕ .

```
1 # Function to get beta shape parameters from moments
2 shape_from_stats <- function(mu = mu.global, sigma = sigma.global) {
3   a <- (mu^2 - mu^3 - mu * sigma^2) / sigma^2
4   b <- (mu - 2 * mu^2 + mu^3 - sigma^2 + mu*sigma^2) / sigma^2
5   shape_ps <- c(a, b)
6   return(shape_ps)
7 }
8
9 # get parameters for distribution of population multiplier, 1/p
10 shapes = shape_from_stats(.163, .012)
11
12 # check prior on p using simulated data from beta distribution
13 x = seq(0, 1, .001)
14 p = dbeta(x, shapes[1], shapes[2])
15 plot(x, p, typ = "l", xlim = c(0, 1))
```





Diagram the Bayesian network

1. Develop a hierarchical Bayesian model (also called a state space model) of the lynx population in the management unit. Diagram the Bayesian network (the DAG) of knowns and unknowns and write out the posterior and factored joint distribution. Use a lognormal distribution to model the true lynx population size over time. Use a Poisson distribution for the data model relating the true, unobserved state (the total population size) to the observed data (number of family groups).



2. An alternative approach, which is slightly more difficult to code, is to model the process as:

$$\text{negative binomial}(N_t \mid \lambda(N_{t-1} - H_{t-1}, \rho))$$

,

and model the data as:

$$\text{binomial}(y_t \mid \text{round}(N_t \phi), p)$$

,

where p is a detection probability. Explain why this second formulation might be better than the formulation you are using. (It turns out they give virtually identical results.)

ANSWER


Fitting the Model

Now you'll estimate the marginal posterior distribution of the unobserved, true state over time (\mathbf{N}), the parameters in the model λ and ϕ as well as the process variance and observation variance. You'll also summarize the marginal posterior distributions of the parameters and unobserved states. A note about the data. Each row in the data file gives the observed number of family groups for that year in column 2 and that year's harvest in column 3. The harvest in each row influences the population size in the next row. So, for example, the 2016 harvest influences the 2017 population size.

Before you begin it's very helpful to use simulated data to verify initial values and model. We simulate the true state by choosing some biologically reasonable values for model parameters and "eyeballing" the fit of the true state to the data. You can then use these simulated values for initial conditions (see the `inits` list below). This is of particular importance because failing to give reasonable initial conditions for dynamic models can cause problems in model fitting. Remember, supply initial conditions for *all* unobserved quantities in the posterior distribution (even those that do not have priors).

```
1 y <- LynxFamilies
2 endyr <- nrow(y)
3 n <- numeric(endyr + 1)
4 mu <- numeric(endyr + 1)
5 fg <- numeric(endyr + 1)
6 phi <- 0.16
7 lambda <- 1.07
8 sigma.p <- 0.2
9
10 n[1] <- y$census[1] / phi # n in the unit of individuals
11 mu[1] <- n[1] # mean from deterministic model to simulate
12 fg[1] <- n[1] * phi # Nt in the unit of
13
14 for (t in 2:(endyr + 1)) {
15   mu[t] <- lambda * (n[t - 1] - y$harvest[t - 1])
16   n[t] <- rlnorm(1, log(mu[t]), sigma.p)
17   fg[t] <- n[t] * phi
18 }
19
20 plot(y$year, y$census, ylim = c(0, 100), xlab = "Year", ylab = "Family group",
21      main = "Simulated data")
22 lines(y$year, fg[1:length(y$year)])
```

Simulated data




```

1  ## visually match simulated data with observations for initial conditions
2  endyr = nrow(y)
3  n = numeric(endyr + 1)
4  mu = numeric(endyr + 1) #use this for family groups
5  lambda = 1.1
6  sigma.p = .00001
7  n[1] = y$census[1]
8
9  for(t in 2:(endyr + 1)) {
10     n[t] <- lambda * (y$census[t - 1] - .16 * y$harvest[t - 1]) # use this for family groups
11 }
12
13 plot(y$year, y$census, ylim = c(0, 100), xlab = "Year", ylab = "Family group",
14      main = "Simulated data")
15 lines(y$year, n[1:length(y$year)])

```

Simulated data



Here's your starting code:

```

1 data = list(
2   y.endyr = endyr,
3   y.a = shapes[1],
4   y.b = shapes[2],
5   y.H = y$harvest,
6   y = y$census)
7
8 inits = list(
9   list(lambda = 1.2, sigma.p = .01, N = n),
10  list(lambda = 1.01, sigma.p = .2, N = n * 1.2),
11  list(lambda = .95, sigma.p = .5, N = n * .5))

```

1. Write the JAGS model to estimate the marginal posterior distribution of the unobserved, true state over time (N), the parameters in the model λ and ϕ as well as the process variance and observation variance. Include a summary the marginal posterior distributions of the parameters and unobserved states.

```

1 {
2   sink("lynxmodel.R")
3   cat("
4   model{
5
6

```

```

7  # Priors
8
9  phi ~ dbeta(y.a, y.b)
10 lambda ~ dgamma(.5, .001)
11 sigma.p ~ dunif(0,100)
12 tau.p <- 1/sigma.p^2
13 fg[1] ~ dpois(y[1])
14 N[1] ~ dlnorm(log(y[1] / phi), tau.p)
15
16
17 # Likelihood
18
19 for(t in 2:y.endyr){
20   # Data
21   y[t] ~ dpois(phi * N[t])
22 }
23
24
25 # Process
26
27 for(t in 2:(y.endyr+1)){
28   log.mu[t] <- log(max(0.0001,lambda * (N[t-1] - y.H[t-1])))
29   N[t] ~ dlnorm(log.mu[t], tau.p)
30   fg[t] <- N[t] * phi
31 }
32
33
34 # Bayesian p values
35
36 for(t in 1:y.endyr){
37   y.sim[t] ~ dpois(phi * N[t])
38   sq.data[t] <- (y[t] - (phi*N[t]))^2
39   sq.sim[t] <- (y.sim[t] - (phi*N[t]))^2
40 }
41
42 # Predictive check for discrepancy
43 dis.data <- sum(sq.data)
44 dis.sim <- sum(sq.sim)
45 p.dis <- step(dis.sim - dis.data)    # calculate Bayesian test statistic
46
47 # Derived quantity to examine autocorrelation in residuals
48 for(t in 2:y.endyr){
49   e[t] <- y[t] - N[t]
50 }
51
52 }
53   ", fill = TRUE)
54 sink()
55 }

```

```

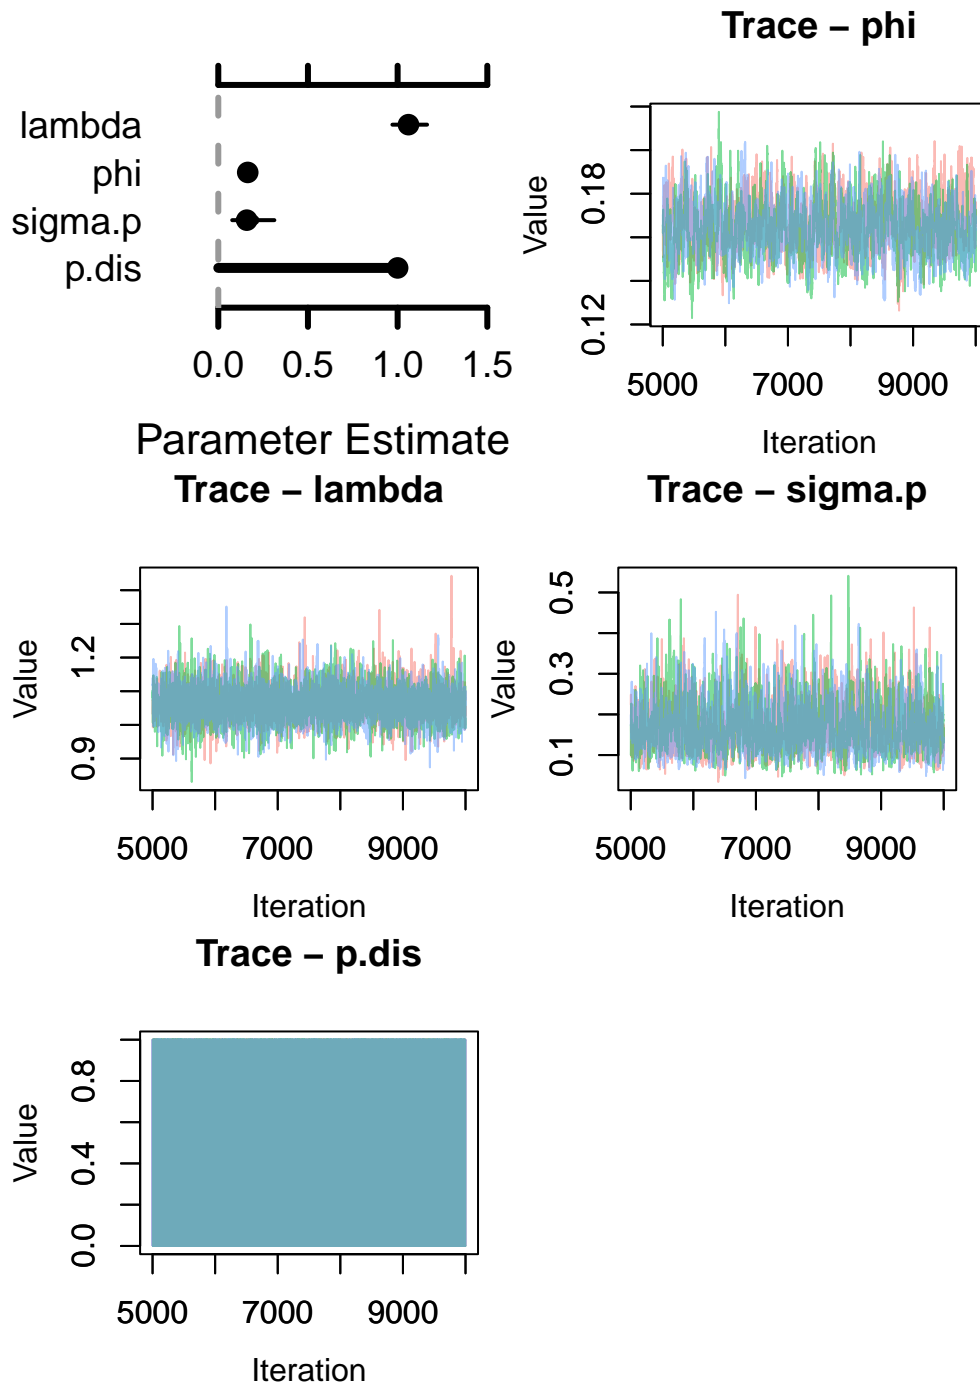
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:

```

```
## Observed stochastic nodes: 18
## Unobserved stochastic nodes: 43
## Total graph size: 309
##
## Initializing model
```

2. Check MCMC chains for model parameters, process variance, and latent states for convergence. This will probably require using the `excl` option in `MCMCsummary`.

##		mean	sd	2.5%	50%	97.5%	Rhat	n.eff
##	N[1]	498.5102	78.8302	369.8333	490.0456	679.2531	1.00	2587
##	N[2]	471.0153	56.8983	370.7223	467.1786	594.2512	1.01	1928
##	N[3]	390.5288	44.8662	309.5329	387.9613	485.2600	1.01	2010
##	N[4]	314.1365	38.8515	242.9704	312.2157	395.6017	1.01	1974
##	N[5]	272.7012	34.4648	209.1106	271.2013	344.6715	1.01	1963
##	N[6]	243.5696	32.3595	184.5179	241.8920	311.7830	1.01	2117
##	N[7]	229.3634	31.2023	172.3762	227.8653	294.3737	1.00	2009
##	N[8]	248.1195	32.5828	190.1056	245.9645	318.6538	1.00	2041
##	N[9]	266.4227	33.8216	206.7627	264.5780	338.3328	1.00	2015
##	N[10]	273.6805	34.0243	211.9451	271.9310	344.6694	1.00	2057
##	N[11]	303.7044	37.2782	237.8119	301.1415	384.2679	1.00	2123
##	N[12]	299.8663	36.1953	237.3985	297.0393	379.2211	1.00	2182
##	N[13]	241.2155	30.1371	186.5551	239.7199	304.5715	1.00	2109
##	N[14]	220.4485	27.9768	171.8551	218.4779	282.0419	1.01	2041
##	N[15]	165.3553	24.2740	121.8874	164.0078	216.7763	1.01	1948
##	N[16]	136.1992	22.2506	94.6656	135.5919	181.3177	1.01	2000
##	N[17]	146.0732	22.5008	104.2602	145.2400	192.8278	1.01	2018
##	N[18]	184.6193	26.6441	137.8896	182.6077	242.9341	1.00	2298
##	N[19]	207.2283	33.4968	149.3377	204.2388	281.2687	1.00	3073
##	N[20]	182.7110	52.5947	104.1715	174.3972	308.3178	1.00	5553
##	e[2]	-387.0153	56.8983	-510.2512	-383.1786	-286.7223	1.01	1928
##	e[3]	-327.5288	44.8662	-422.2600	-324.9613	-246.5329	1.01	2010
##	e[4]	-265.1365	38.8515	-346.6017	-263.2157	-193.9704	1.01	1974
##	e[5]	-228.7012	34.4648	-300.6715	-227.2013	-165.1106	1.01	1963
##	e[6]	-204.5696	32.3595	-272.7830	-202.8920	-145.5179	1.01	2117
##	e[7]	-197.3634	31.2023	-262.3737	-195.8653	-140.3762	1.00	2009
##	e[8]	-206.1195	32.5828	-276.6538	-203.9645	-148.1056	1.00	2041
##	e[9]	-222.4227	33.8216	-294.3328	-220.5780	-162.7627	1.00	2015
##	e[10]	-231.6805	34.0243	-302.6694	-229.9310	-169.9451	1.00	2057
##	e[11]	-250.7044	37.2782	-331.2679	-248.1415	-184.8119	1.00	2123
##	e[12]	-246.8663	36.1953	-326.2211	-244.0393	-184.3985	1.00	2182
##	e[13]	-206.2155	30.1371	-269.5715	-204.7199	-151.5551	1.00	2109
##	e[14]	-181.4485	27.9768	-243.0419	-179.4779	-132.8551	1.01	2041
##	e[15]	-134.3553	24.2740	-185.7763	-133.0078	-90.8874	1.01	1948
##	e[16]	-122.1992	22.2506	-167.3177	-121.5919	-80.6656	1.01	2000
##	e[17]	-126.0732	22.5008	-172.8278	-125.2400	-84.2602	1.01	2018
##	e[18]	-151.6193	26.6441	-209.9341	-149.6077	-104.8896	1.00	2298
##	e[19]	-169.2283	33.4968	-243.2687	-166.2388	-111.3377	1.00	3073
##	lambda	1.0623	0.0478	0.9708	1.0604	1.1642	1.00	9658
##	p.dis	0.5814	0.4933	0.0000	1.0000	1.0000	1.00	18461
##	phi	0.1645	0.0121	0.1417	0.1642	0.1892	1.01	1014
##	sigma.p	0.1684	0.0615	0.0775	0.1590	0.3127	1.00	1756



3. Conduct posterior predictive checks by simulating a new dataset for family groups (f_t) at every MCMC iteration. Calculate a Bayesian p value using the sums of squared discrepancy between the observed and the predicted number of family groups based on observed and simulated data,

$$T^{observed} = \sum_{t=1}^n (f_t^{observed} - N_t \phi)^2$$

$$T^{model} = \sum_{t=1}^n (f_t^{simulated} - N_t \phi)^2$$

The Bayesian p value is the proportion of MCMC iterations for which $T_{model} > T_{obs}$.

Assure yourself that the process model adequately accounts for temporal autocorrelation in the residuals—allowing the assumption that they are independent and identically distributed. To do this, include a derived quantity

$$e_t = y_t - N_t \phi$$


in your JAGS code and coda object. Use the following code or something like it to examine how autocorrelation in the residuals changes with time lag.

```
1 acf(unlist(MCMCpstr(z, param = "e", func = mean)), main = "", lwd = 3, ci = 0)
```

4. Write a paragraph describing how to interpret the plot produced by this function.

ANSWER

5. Plot the median of the marginal posterior distribution of the number of lynx family groups over time (1998-2016) including a highest posterior density interval. Include your forecast for 2017 (the predictive process distribution) in this plot.



Code

```
1 knitr::opts_chunk$set(  
2   echo = FALSE,  
3   message = FALSE,  
4   warning = FALSE,  
5   attr.source = ".numberLines"  
6 )  
7 library(BayesNSF)  
8 library(rjags)  
9 library(MCMCvis)  
10 library(HDInterval)  
11 set.seed(10)  
12 library(BayesNSF)  
13 library(rjags)  
14 library(MCMCvis)  
15 library(HDInterval)  
16 set.seed(10)  
17 # Function to get beta shape parameters from moments  
18 shape_from_stats <- function(mu = mu.global, sigma = sigma.global) {  
19   a <- (mu^2 - mu^3 - mu * sigma^2) / sigma^2  
20   b <- (mu - 2 * mu^2 + mu^3 - sigma^2 + mu*sigma^2) / sigma^2  
21   shape_ps <- c(a, b)  
22   return(shape_ps)  
23 }  
24  
25 # get parameters for distribution of population multiplier, 1/p  
26 shapes = shape_from_stats(.163, .012)  
27  
28 # check prior on p using simulated data from beta distribution  
29 x = seq(0, 1, .001)  
30 p = dbeta(x, shapes[1], shapes[2])  
31 plot(x, p, typ = "l", xlim = c(0, 1))  
32 DiagrammeR::grViz("  
33   digraph mrdag {  
34     graph [rankdir=TB, layout=neato]  
35  
36     node [shape=plaintext, height=0.3, width=0.3]  
37     Y      [label=<<I>Y@_{t}</I>>, pos='3,1!']  
38  
39     N      [label=<<I>N@_{t}</I>>, pos='2,0!']  
40  
41     phi    [label='&phi;', pos='4,-2!']  
42  
43     H      [label=<<I>H</I>>, pos='3,-1!']  
44  
45     lambda [label='&lambda;', pos='2,-1!']  
46  
47     sigma  [label='<I>&sigma;@_{p-t}^2</I>', pos='1,-1!']  
48  
49     Nt1    [label=<<I>N@_{t-1}</I>>, pos='0,-1!']  
50  
51     N1     [label=<<I>N@_{1}</I>>, pos='-0,-2!']
```

```

52     edge [arrowhead='vee']
53     N -> Y
54     phi -> Y
55     sigma -> N
56     lambda -> N
57     Nt1 -> N
58
59
60     H -> N      [style=dashed];
61     N1 -> Nt1   [style=dotted];
62     phi -> N1   [style=dashed];
63     }
64     ", height = 190)
65 y <- LynxFamilies
66 endyr <- nrow(y)
67 n <- numeric(endyr + 1)
68 mu <- numeric(endyr + 1)
69 fg <- numeric(endyr + 1)
70 phi <- 0.16
71 lambda <- 1.07
72 sigma.p <- 0.2
73
74 n[1] <- y$census[1] / phi # n in the unit of individuals
75 mu[1] <- n[1] # mean from deterministic model to simulate
76 fg[1] <- n[1] * phi # Nt in the unit of
77
78 for (t in 2:(endyr + 1)) {
79     mu[t] <- lambda * (n[t - 1] - y$harvest[t - 1])
80     n[t] <- rlnorm(1, log(mu[t]), sigma.p)
81     fg[t] <- n[t] * phi
82 }
83
84 plot(y$year, y$census, ylim = c(0, 100), xlab = "Year", ylab = "Family group",
85      main = "Simulated data")
86 lines(y$year, fg[1:length(y$year)])
87 ## visually match simulated data with observations for initial conditions
88 endyr = nrow(y)
89 n = numeric(endyr + 1)
90 mu = numeric(endyr + 1) #use this for family groups
91 lambda = 1.1
92 sigma.p = .00001
93 n[1] = y$census[1]
94
95 for(t in 2:(endyr + 1)) {
96     n[t] <- lambda * (y$census[t - 1] - .16 * y$harvest[t - 1]) # use this for family groups
97 }
98
99 plot(y$year, y$census, ylim = c(0, 100), xlab = "Year", ylab = "Family group",
100      main = "Simulated data")
101 lines(y$year, n[1:length(y$year)])
102 data = list(
103     y.endyr = endyr,
104     y.a = shapes[1],
105     y.b = shapes[2],

```

```

106     y.H = y$harvest,
107     y = y$census)
108
109 inits = list(
110     list(lambda = 1.2, sigma.p = .01, N = n),
111     list(lambda = 1.01, sigma.p = .2, N = n * 1.2),
112     list(lambda = .95, sigma.p = .5, N = n * .5))
113 {
114     sink("lynxmodel.R")
115     cat("
116     model{
117
118
119     # Priors
120
121     phi ~ dbeta(y.a, y.b)
122     lambda ~ dgamma(.5, .001)
123     sigma.p ~ dunif(0,100)
124     tau.p <- 1/sigma.p^2
125     fg[1] ~ dpois(y[1])
126     N[1] ~ dlnorm(log(y[1] / phi), tau.p)
127
128
129     # Likelihood
130
131     for(t in 2:y.endyr){
132     # Data
133         y[t] ~ dpois(phi * N[t])
134     }
135
136
137     # Process
138
139     for(t in 2:(y.endyr+1)){
140         log.mu[t] <- log(max(0.0001, lambda * (N[t-1] - y.H[t-1])))
141         N[t] ~ dlnorm(log.mu[t], tau.p)
142         fg[t] <- N[t] * phi
143     }
144
145
146     # Bayesian p values
147
148     for(t in 1:y.endyr){
149         y.sim[t] ~ dpois(phi * N[t])
150         sq.data[t] <- (y[t] - (phi*N[t]))^2
151         sq.sim[t] <- (y.sim[t] - (phi*N[t]))^2
152     }
153
154     # Predictive check for discrepancy
155     dis.data <- sum(sq.data)
156     dis.sim <- sum(sq.sim)
157     p.dis <- step(dis.sim - dis.data)    # calcualte Bayesian test statistic
158

```

```

159 # Derived quantity to examine autocorrelation in residuals
160 for(t in 2:y.endyr){
161   e[t] <- y[t] - N[t]
162 }
163
164 }
165   ", fill = TRUE)
166   sink()
167 }
168 # Setup MCMC chain
169 n.adapt = 3000
170 n.update = 10000
171 n.iter = 10000
172
173 jm = jags.model("lynxmodel.R",
174               data = data, inits = inits, n.chains = length(inits), n.adapt = n.adapt)
175 update(jm, n.iter = n.update)
176
177 zm = coda.samples(jm, variable.names = c("lambda", "phi", "sigma.p", "N", "p.dis", "e"),
178                 n.iter = n.iter, n.thin = 1)
179 MCMCsummary(zm, round = 4, n.eff = TRUE)
180 # caterpillar plot
181 MCMCplot(zm, params=c("lambda","phi", "sigma.p", "p.dis"))
182
183 # trace plots
184 MCMCtrace(zm, params = 'phi', type = 'trace', pdf = FALSE)
185 MCMCtrace(zm, params = 'lambda', type = 'trace', pdf = FALSE)
186 MCMCtrace(zm, params = 'sigma.p', type = 'trace', pdf = FALSE)
187 MCMCtrace(zm, params = 'p.dis', type = 'trace', pdf = FALSE)
188
189
190 acf(unlist(MCMCpstr(z, param = "e", func = mean)), main = "", lwd = 3, ci = 0)
191 zm1 <- coda.samples(jm, variable.names = c("N","fg"), n.iter = n.iter, n.thin = 1)
192 bound <- MCMCpstr(zm1, params = 'fg', func = function(x) hdi(x,0.95))
193
194 years <- seq(1998,2017,1)
195 plot(y$year, y$census, ylab = "Number of Lynx Family Groups ", xlab = "Years",
196      xlim = c(1998,2017), ylim = c(0,100))
197 lines(years, MCMCpstr(zm1, params = 'fg', func = median)$fg, col="red", lwd = 2)
198 lines(bound$fg[,1] ~ years, lty=2)
199 lines(bound$fg[,2] ~ years, lty=2)
200 abline(h=26, lty=3)
201 abline(h=32, lty=3)
202 # this R markdown chunk generates a code appendix

```